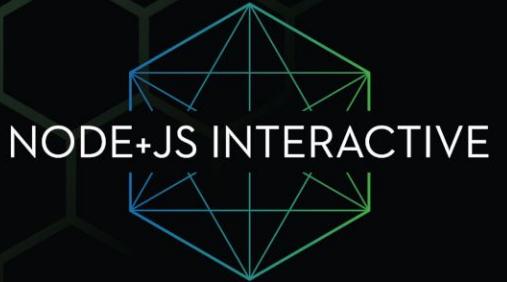


NODE+JS INTERACTIVE

OCTOBER 10-12, 2018 | VANCOUVER, CANADA

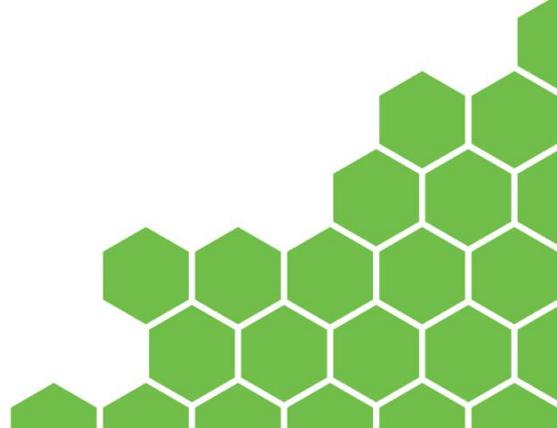


Bringing JavaScript Back to Life

Joyee Cheung, *Igalia*

About Me

- Joyee Cheung / Qiuyi Zhang (张秋怡)
- Hangzhou, China
- Compilers team @ Igalia
- Node.js TSC & Diagnostics WG
- joyeecheung @ GitHub/Twitter

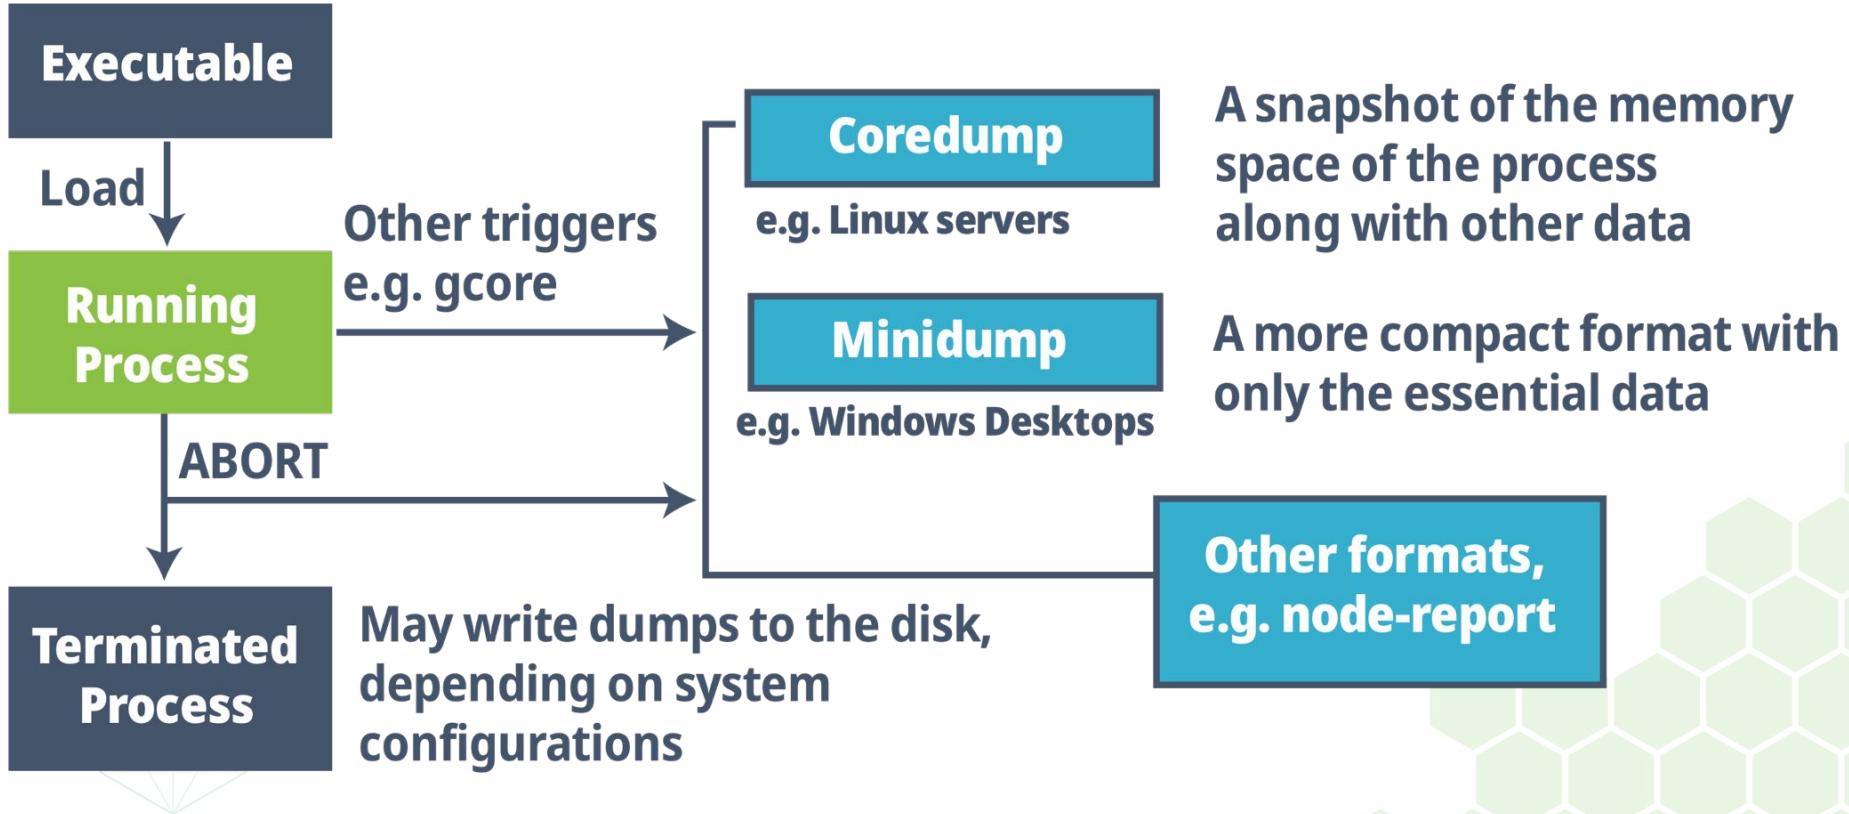


Agenda

- **Introduction to post-mortem diagnosis**
 - Analysis of a coredump
 - Recovering states from dynamic language VMs
- **A tour of llnode**
 - Reconstruct JavaScript values from raw memory
 - Unwinding the Stack
 - JavaScript API of llnode
- **Future of Node/JS Post-mortem Diagnosis**



Introduction to Post-Mortem Diagnostics



Analysis of a Coredump

```
int main(int argc, char* argv[]) {  
    crash(123);  
}
```

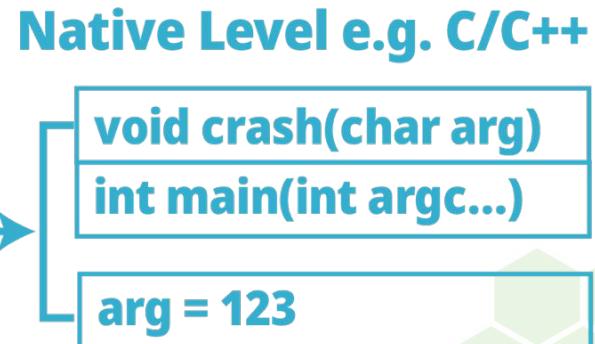


Unstripped ELF
Executable on Linux



Coredump on Linux
(Also in ELF)

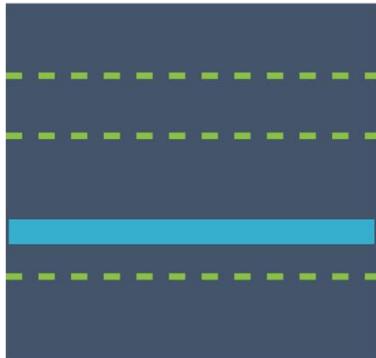
GDB
LLDB
etc



Native Level e.g. C/C++

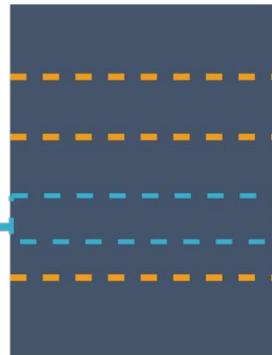
Analysis of a Coredump

```
int main(int argc, char* argv[]) {  
    run_vm(argc, argv);  
}
```



Executable embedding
dynamic language VMs

```
function fn(e) { throw e }  
fn(new Error('crash'))
```



Coredump with dynamic
symbols/type info
generated as data



Dynamic Language Level

Stack Frames ???

Variables ???

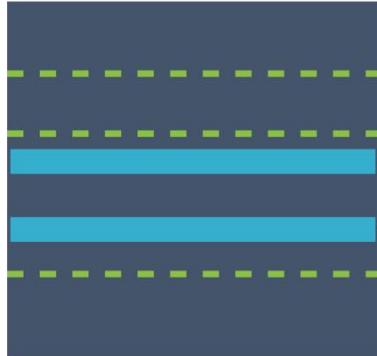
void run_vm(...)
int main(int argc...)

argc = 3

Native Level e.g. C/C++

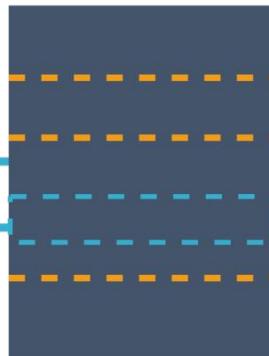
Recovering States from Dynamic Language VMs

```
int main(int argc, char* argv[]) {  
    run_vm(argc, argv);  
}
```



Executable embedding
dynamic language VMs
(+ additional debug info)

```
function fn(e) { throw e }  
fn(new Error('crash'))
```



Coredump with dynamic
symbols/type info
generated as data

Dynamic Language Level

```
function fn(e) {...}  
e = Error { ... }
```

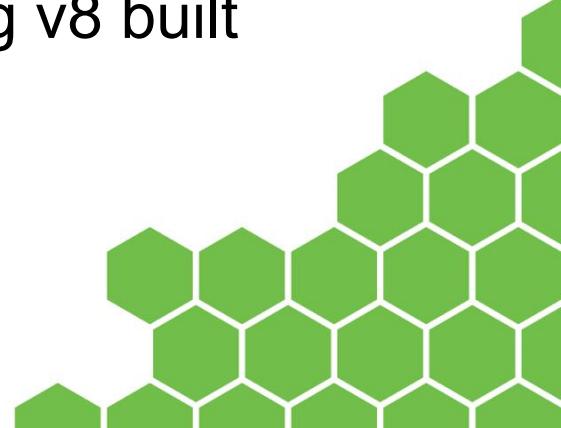
```
void run_vm(...)  
int main(int argc...)
```

```
argc = 3
```

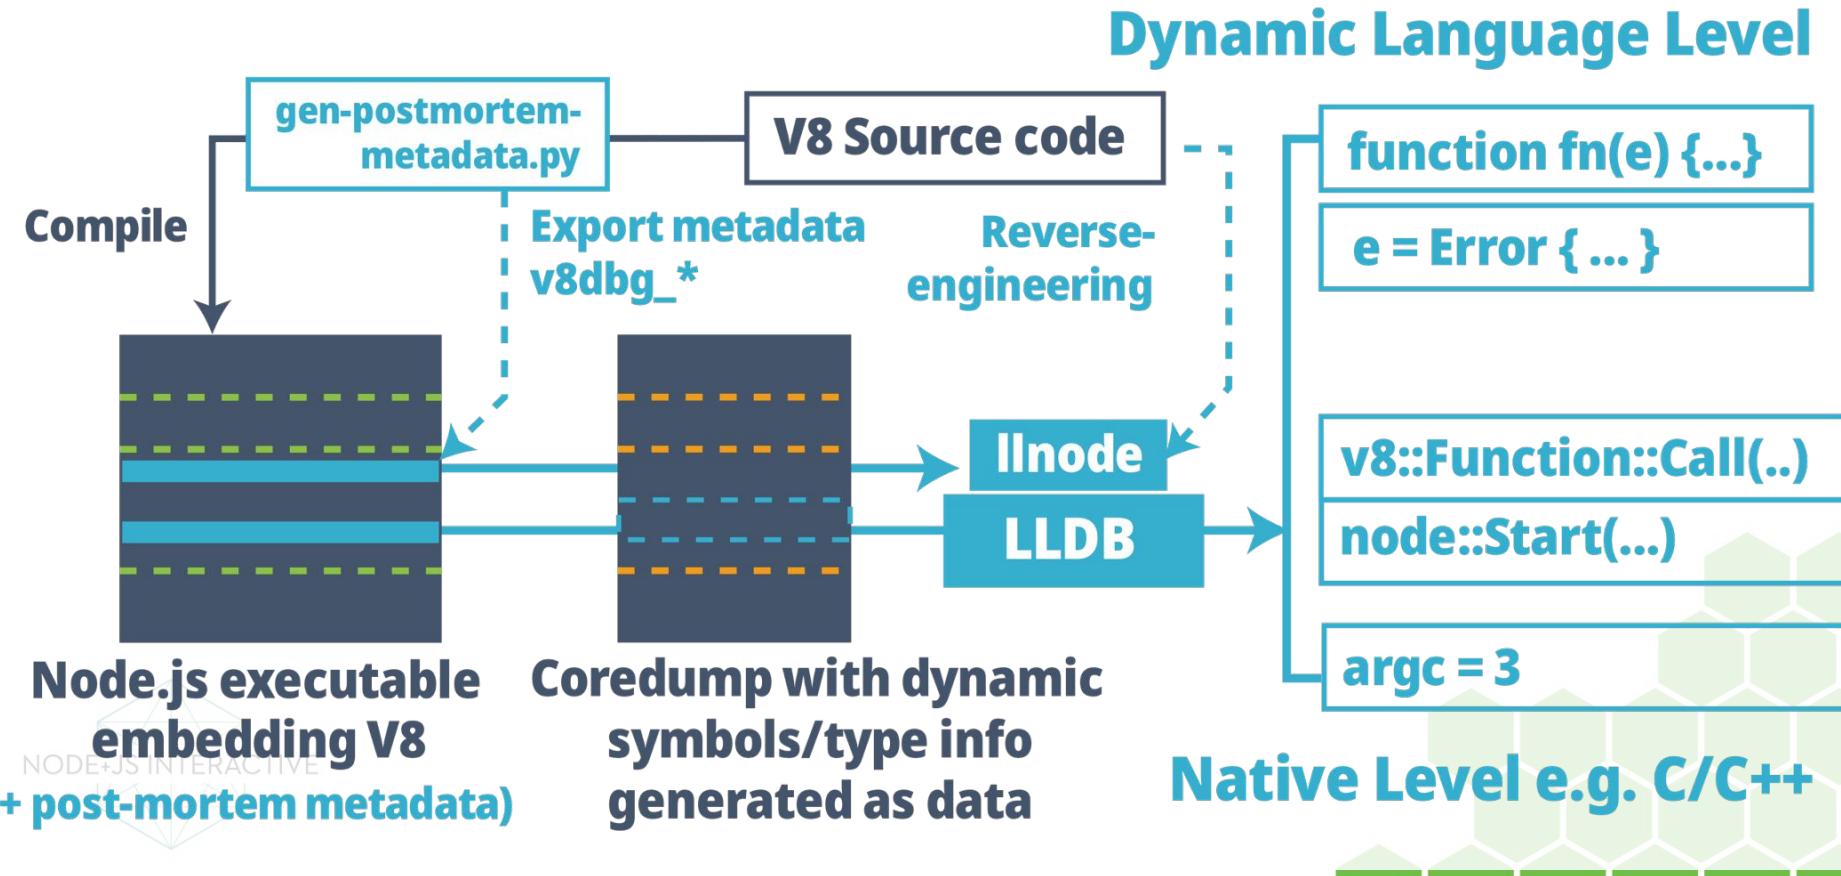
Native Level e.g. C/C++

A Tour of llnode

- Project under the Node.js Diagnostics Working Group
- <https://github.com/nodejs/llnode>
- Plugin of the LLDB debugger
- Can be used to debug both core dumps and live processes of Node.js applications
- Also works with other programs embedding v8 built with post-mortem support
 - e.g. d8 shell

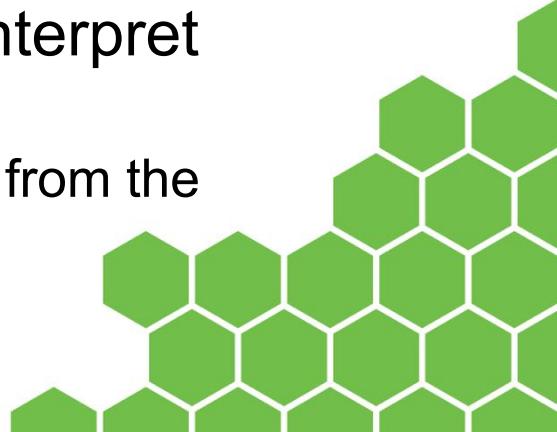
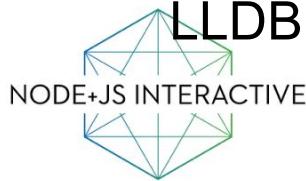


A Tour of llnode

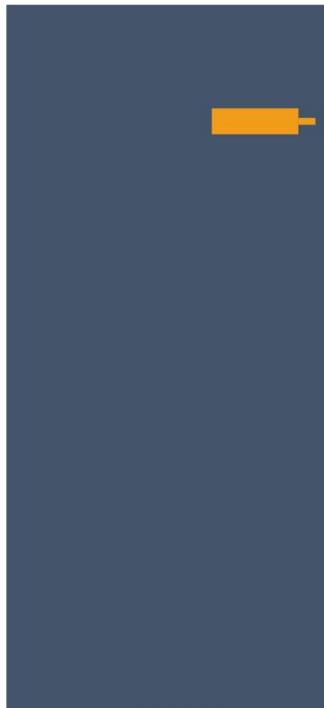


A Tour of llnode

1. (Maintainers) read the V8 source code for reverse-engineering algorithms
2. Load the exact values of `v8dbg_*` offsets/constants from an executable embedding V8 e.g. Node.js
 - If the necessary metadata is not exposed, submit a patch to the upstream
3. Use the algorithm and the metadata to interpret memory blocks in the core dump
 - Infrastructures and cross-platform support come from the LLDB API



Reconstruct JS Values from Raw Memory



Check the last
bit (tag)

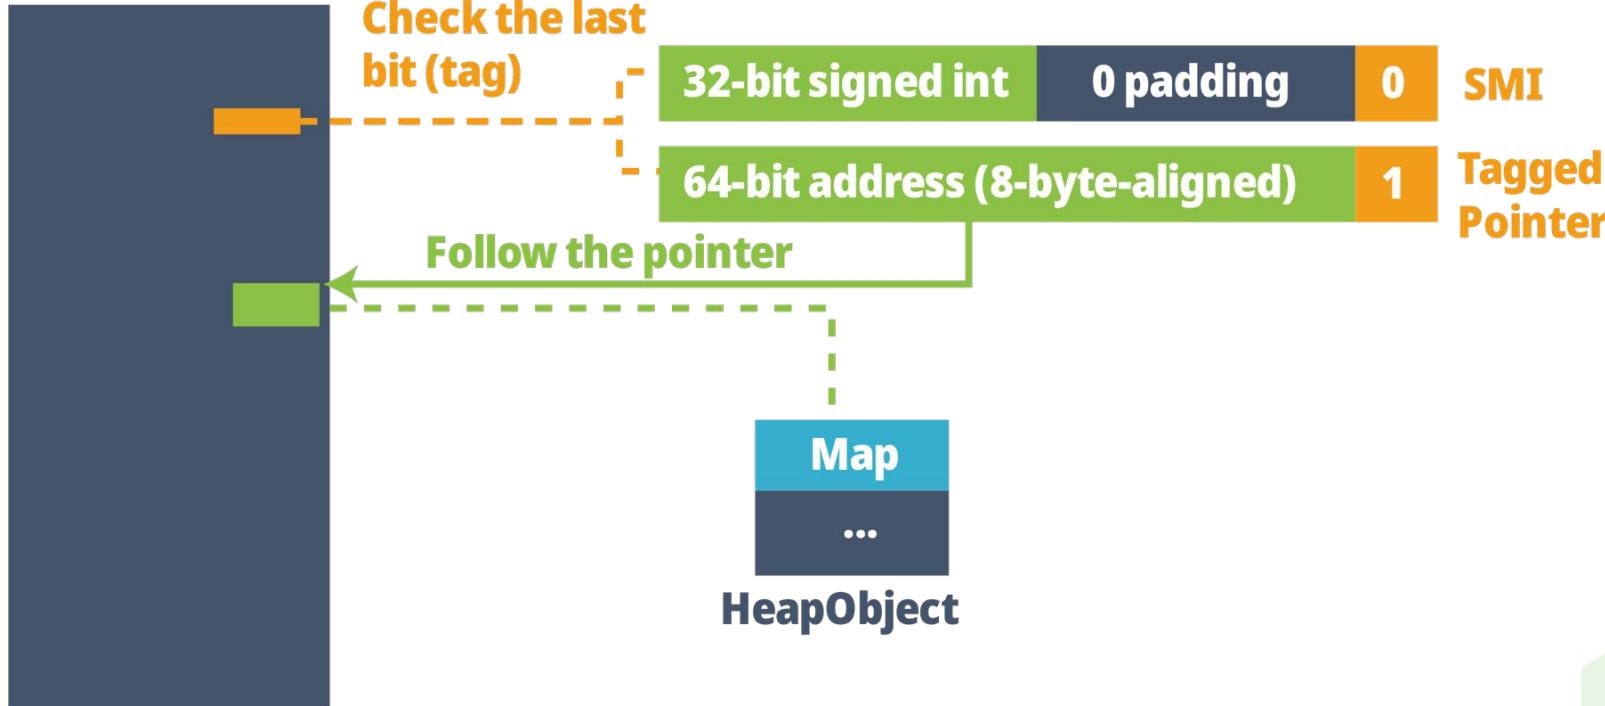


* Layout offsets and
known types are
exposed as metadata

NODEJS INTERACTIVE

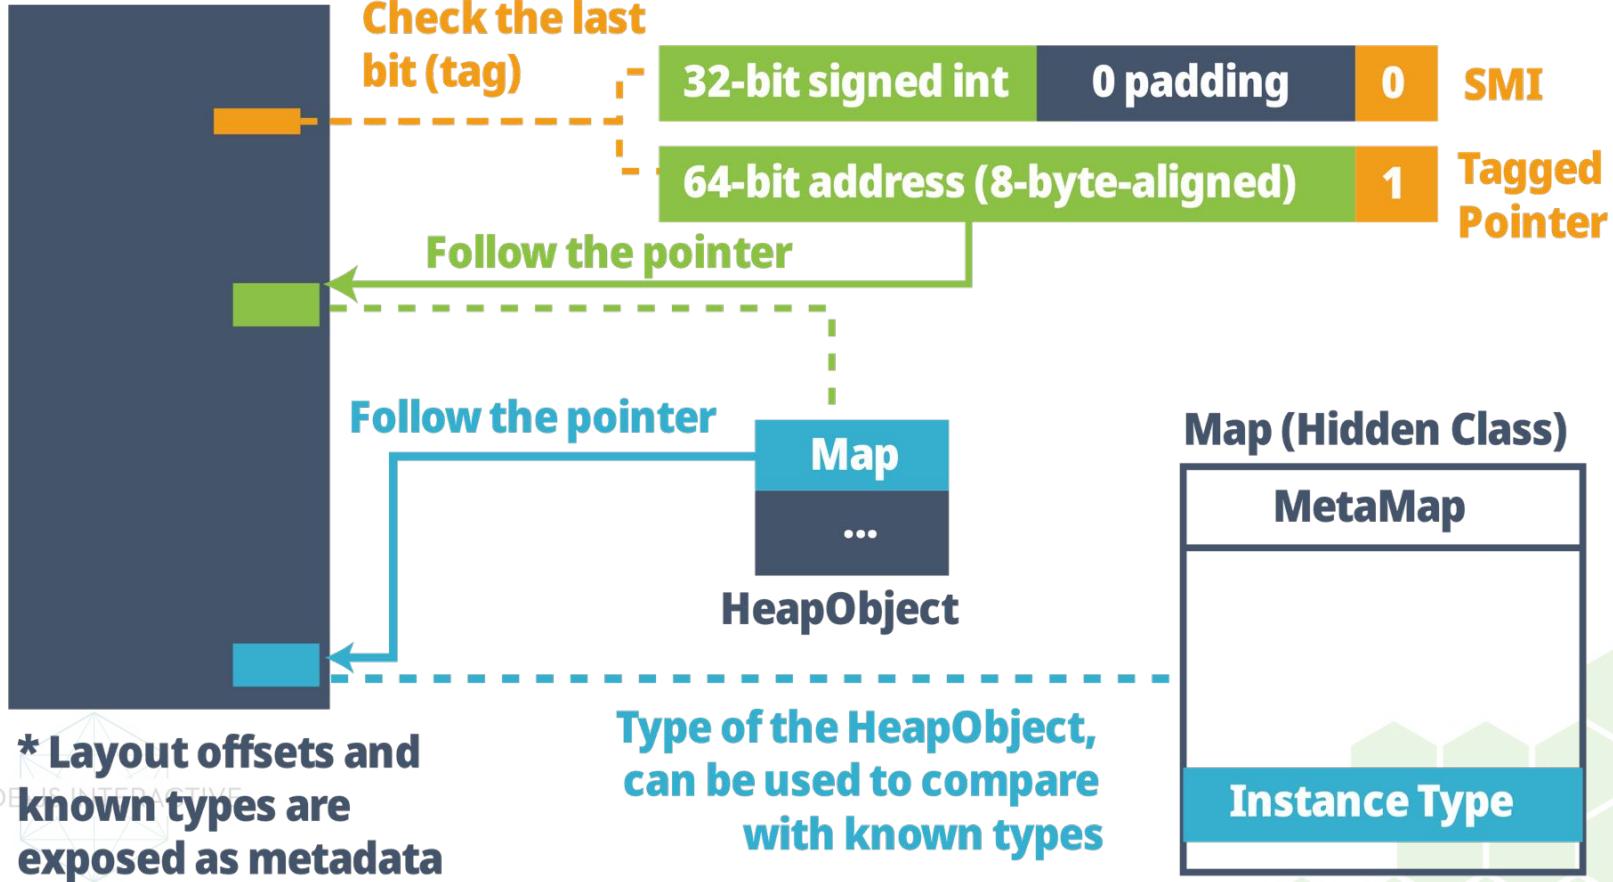


Reconstruct JS Values from Raw Memory



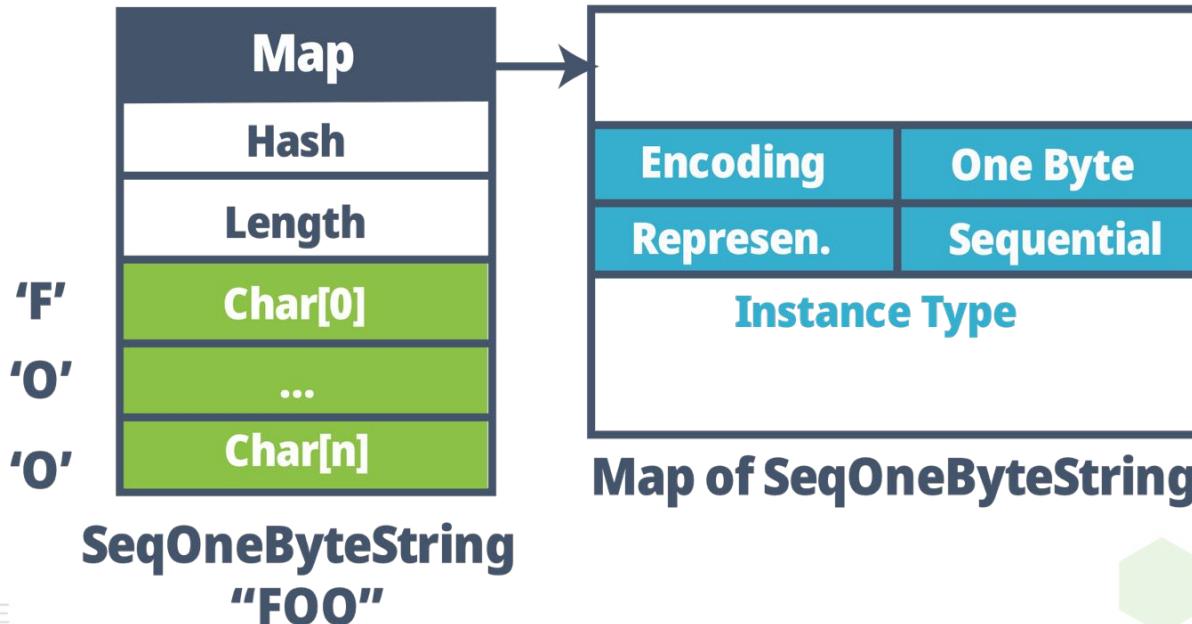
* Layout offsets and known types are exposed as metadata

Reconstruct JS Values from Raw Memory



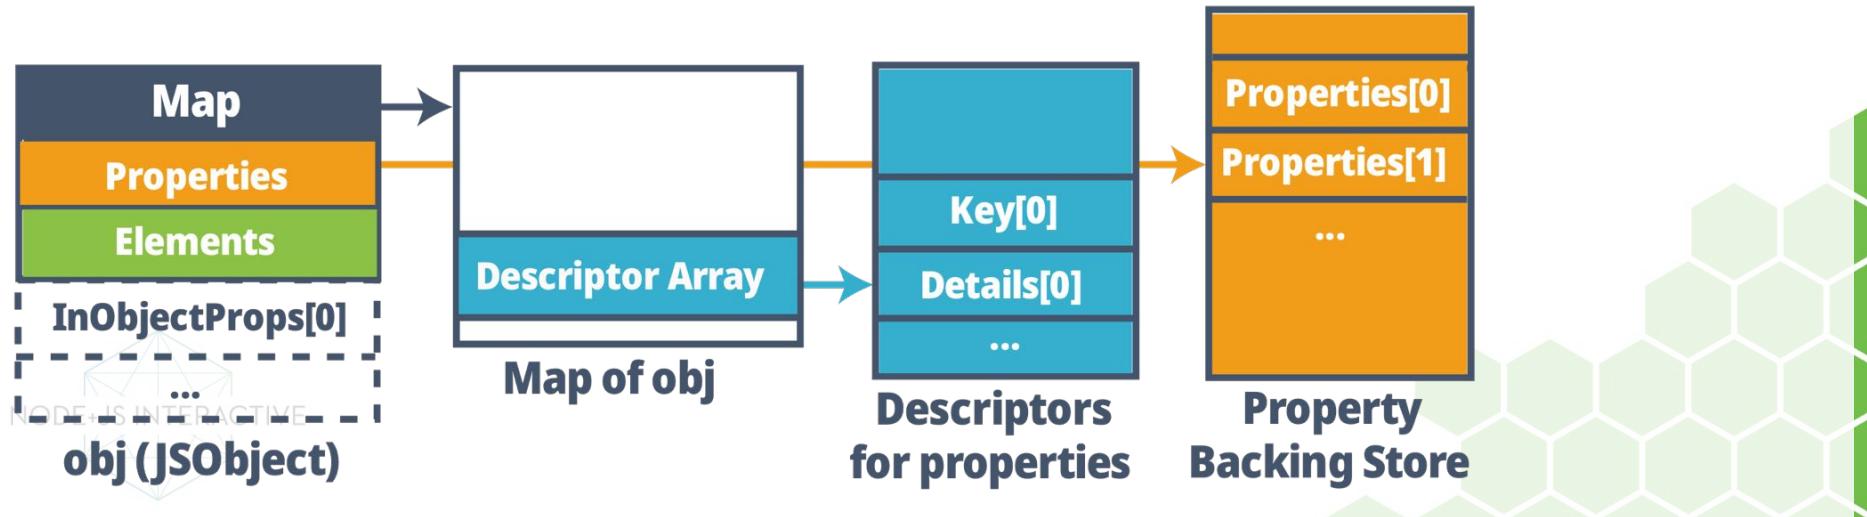
Reconstruct JS Values from Raw Memory

Known Layout



Reconstruct JS Values from Raw Memory

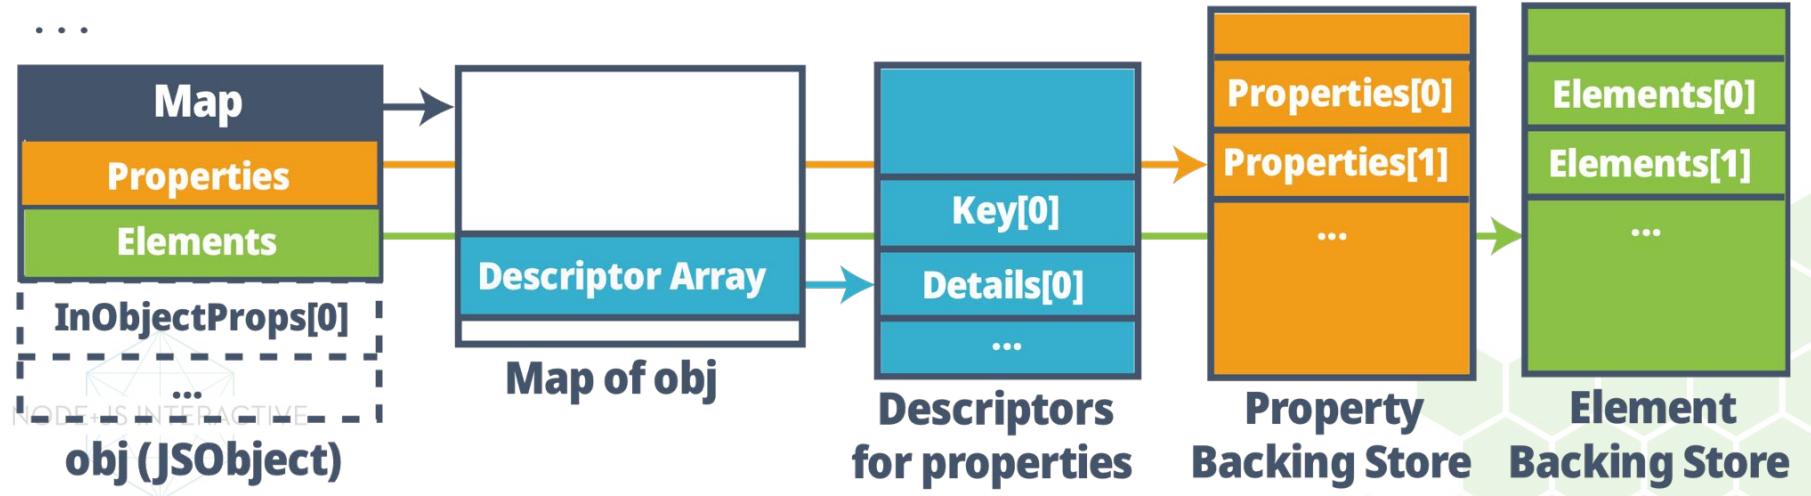
```
const obj = {
  foo: 1, // in-object properties
  ...
};
obj.x = 'test'; // Named properties with descriptors
```



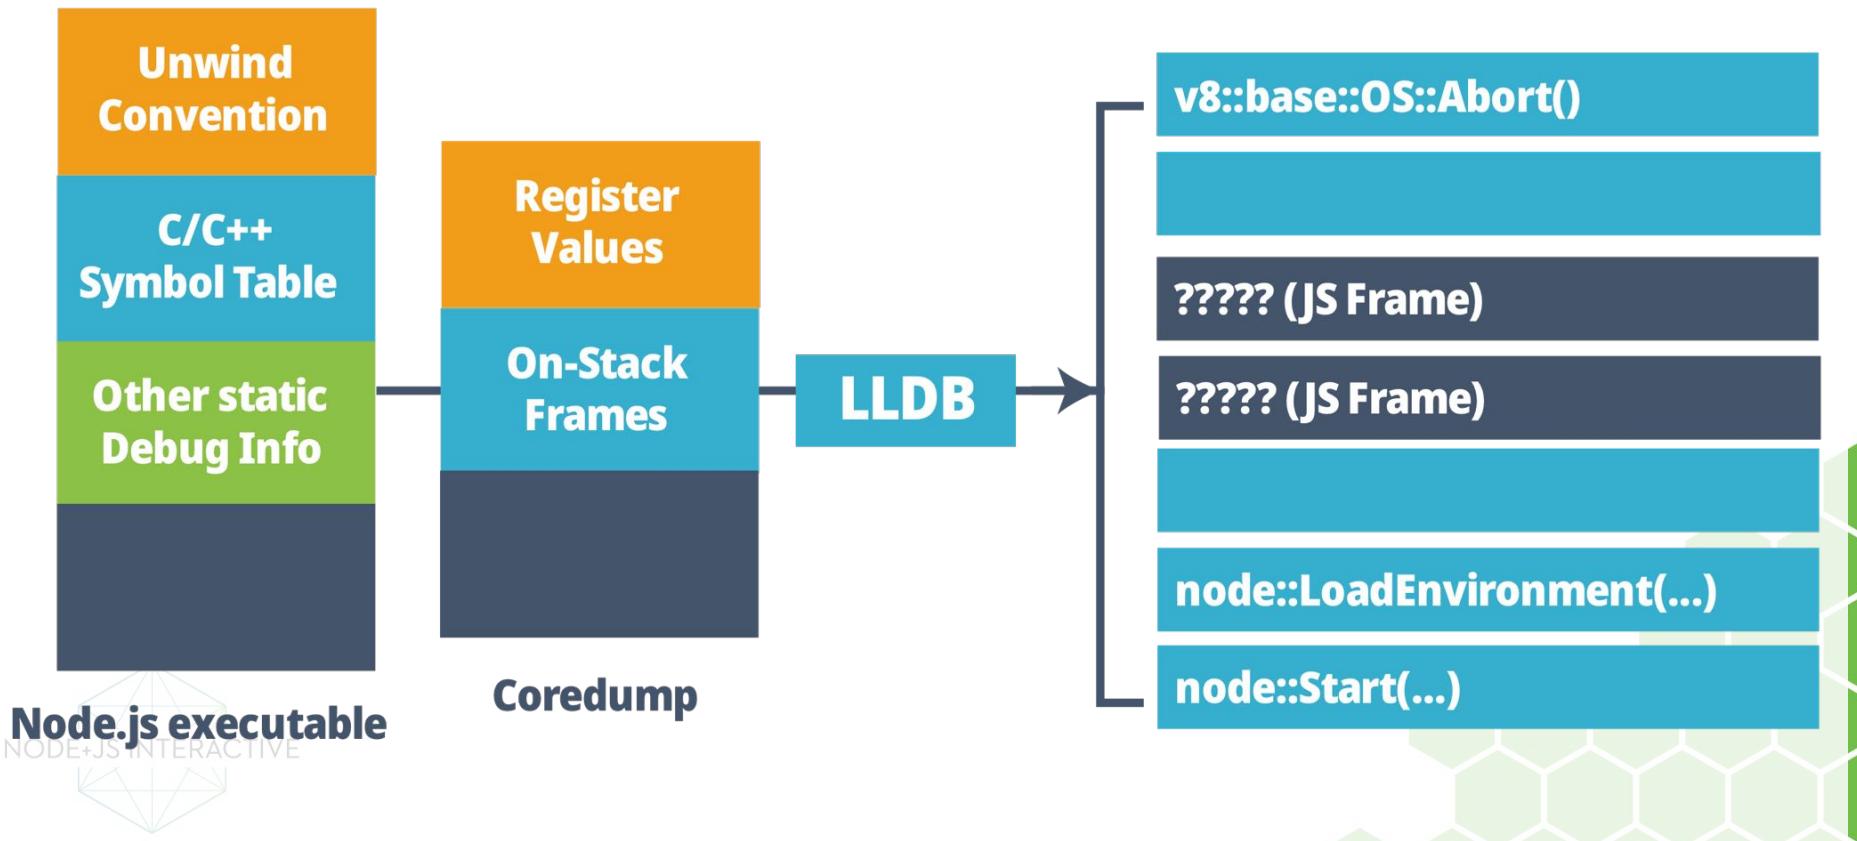
Reconstruct JS Values from Raw Memory

```
const obj = {  
  foo: 1, // in-object properties  
  ...  
};  
obj.x = 'test'; // Named properties with descriptors  
...  
obj[0] = 1; // Indexed elements  
...
```

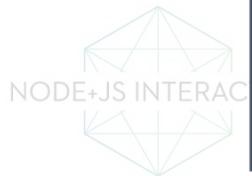
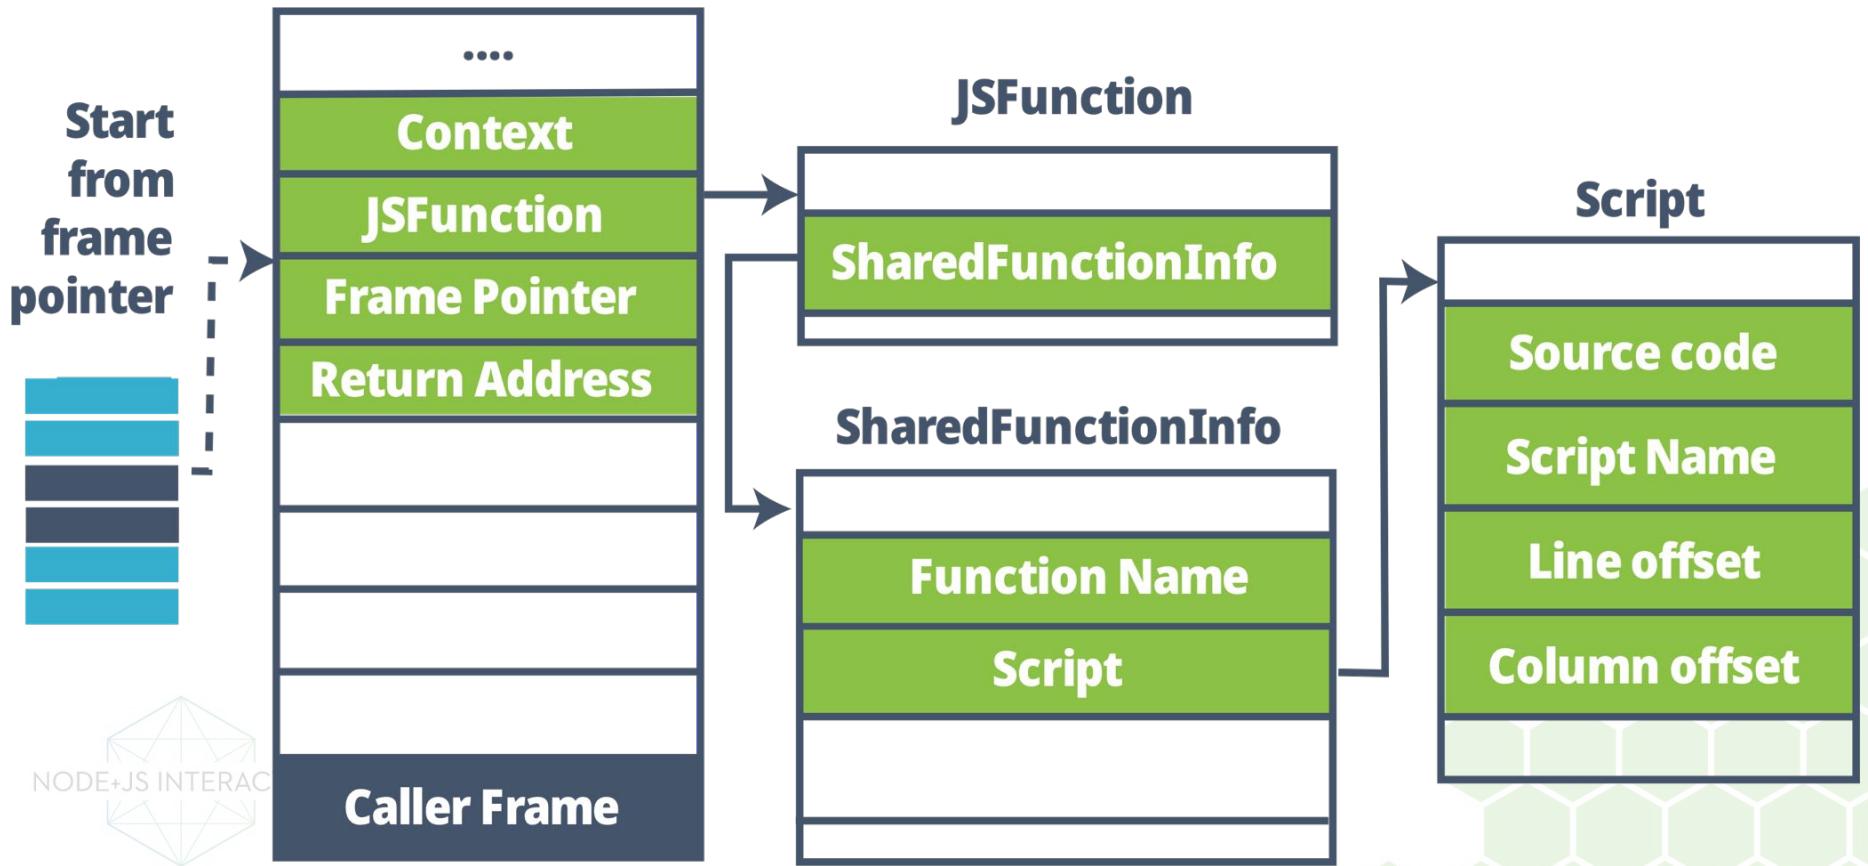
* Objects in dictionary mode have flexible property dictionaries instead of fixed descriptors



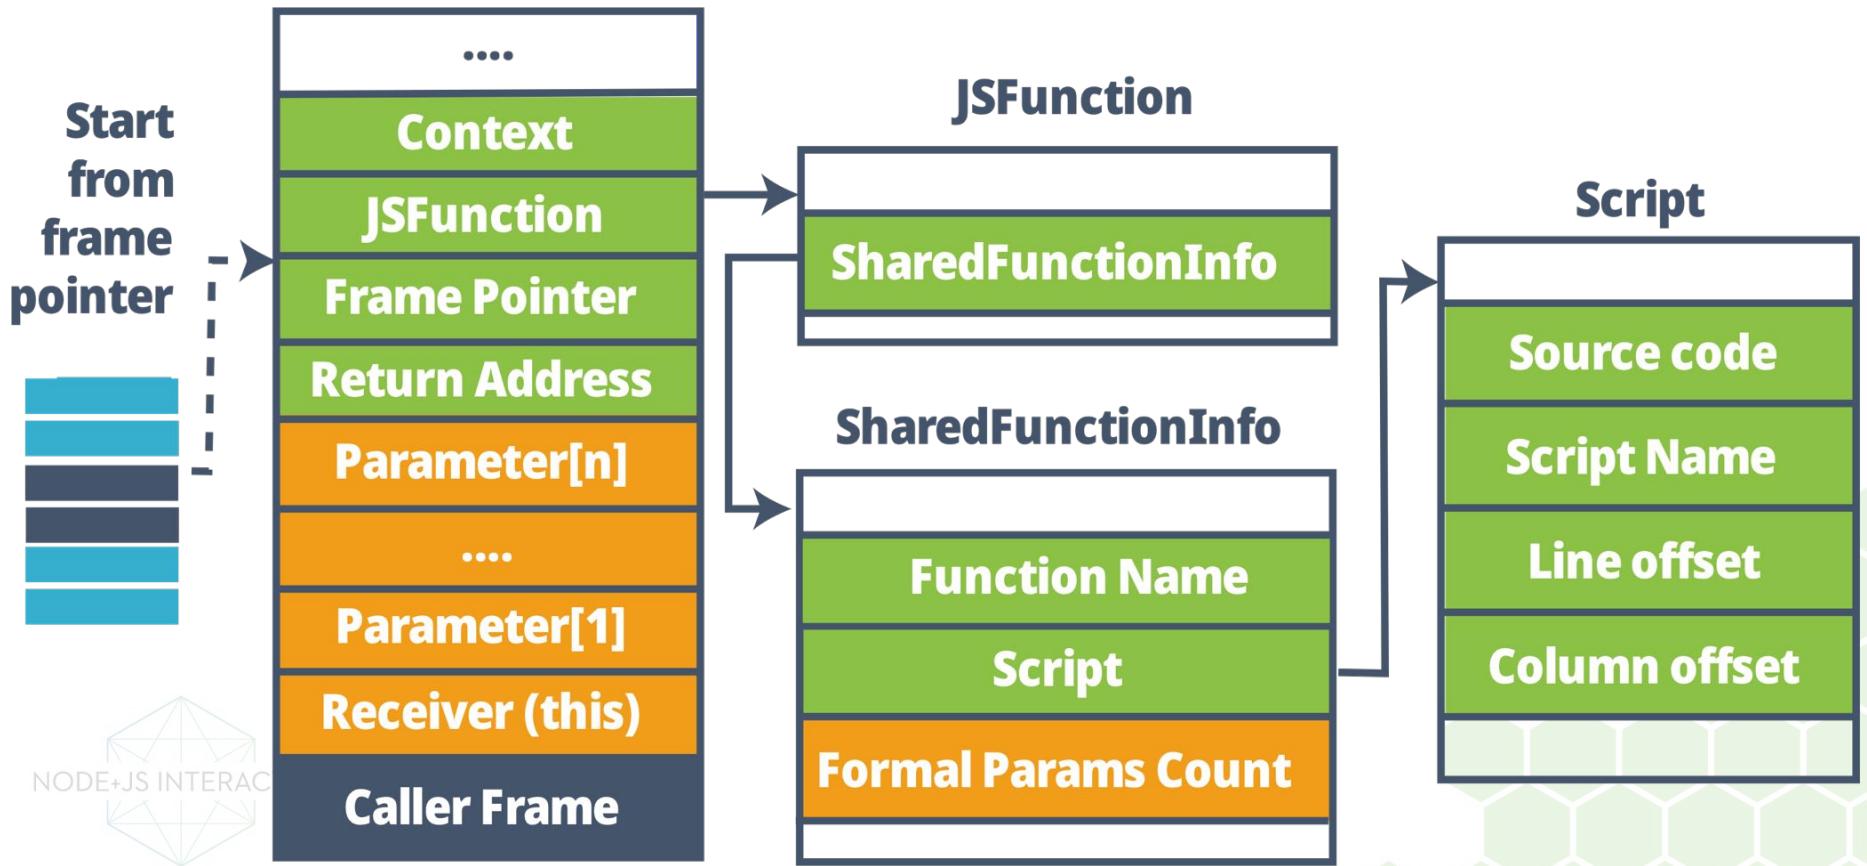
Unwinding the Stack: Native



Unwinding the Stack: JS Symbols



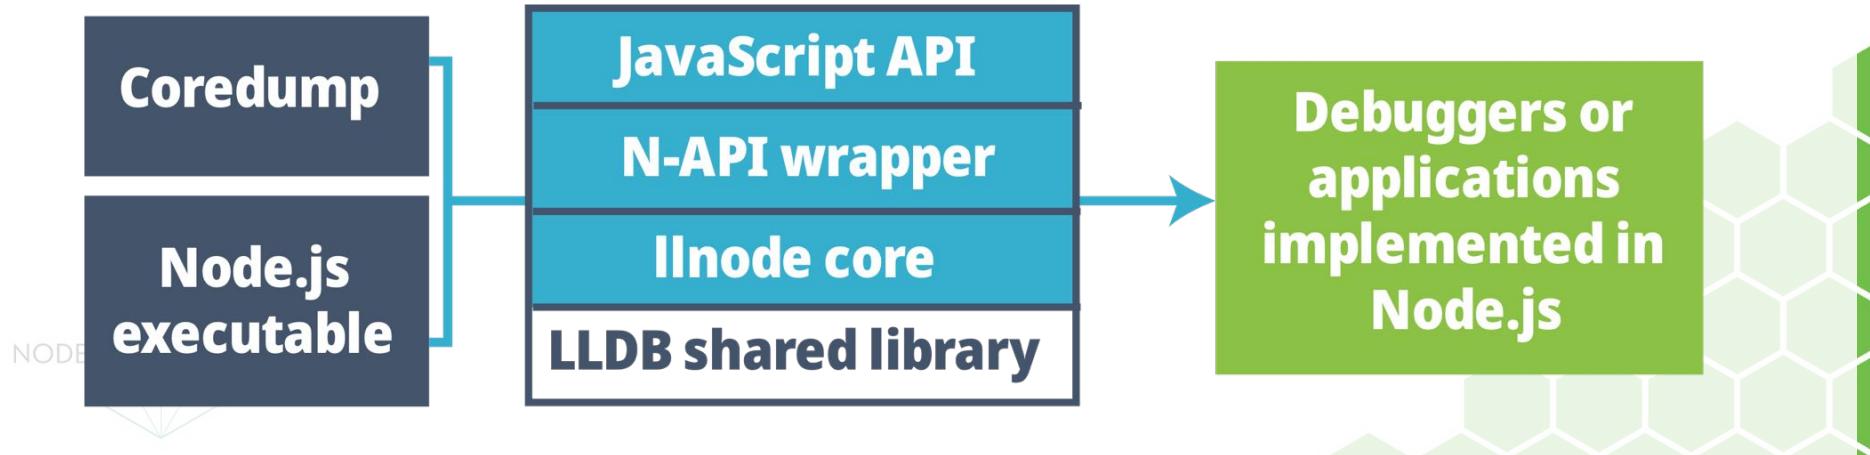
Unwinding the Stack: JS Symbols



JavaScript API of llnode

- As an ordinary Node.js C++ addon (macOS, Linux, FreeBSD, Windows?)
- Primitive interface for now
 - Return strings that are formatted for the LLDB console
 - <https://github.com/nodejs/llnode/blob/master/JSAPI.md>

```
npm install llnode --llnode_build_addon=true
```



JavaScript API of llnode

```
const llnode = require('llnode')
  .fromCoredump('/path/to/core', '/path/to/node');

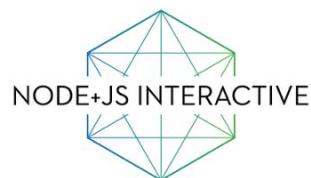
const process = llnode.getProcessObject();

console.log(`Process ${process.pid}: ${process.state}`);

process.threads.forEach((thread) => {
  console.log(`Thread ${thread.threadId}`);
  thread.frames.forEach((frame, index) => {
    console.log(`#${index} ${frame.function}`);
  });
});
```

JavaScript API of llnode

```
llnode.getHeapTypes().forEach((type) => {
  console.log(` ${type.typeName}: ${type.totalSize}`);
  console.log(` ${type.instanceCount} instances`);
  for (const instance of type.instances) {
    console.log(`0x${instance.address}`,
               instance.value);
  }
});
```



The llnode approach

- Highly dependent on the implementation: easy to break whenever V8 changes its **internal** object layout
- Works well enough with the Node.js LTS schedule: only need to support very few versions of V8
- Minimum effort from the VM's side
 - It works!
 - No overhead during runtime
- Fragmentation among llnode, mdb_v8, .etc
 - Limited to the platform supported by the native debugger

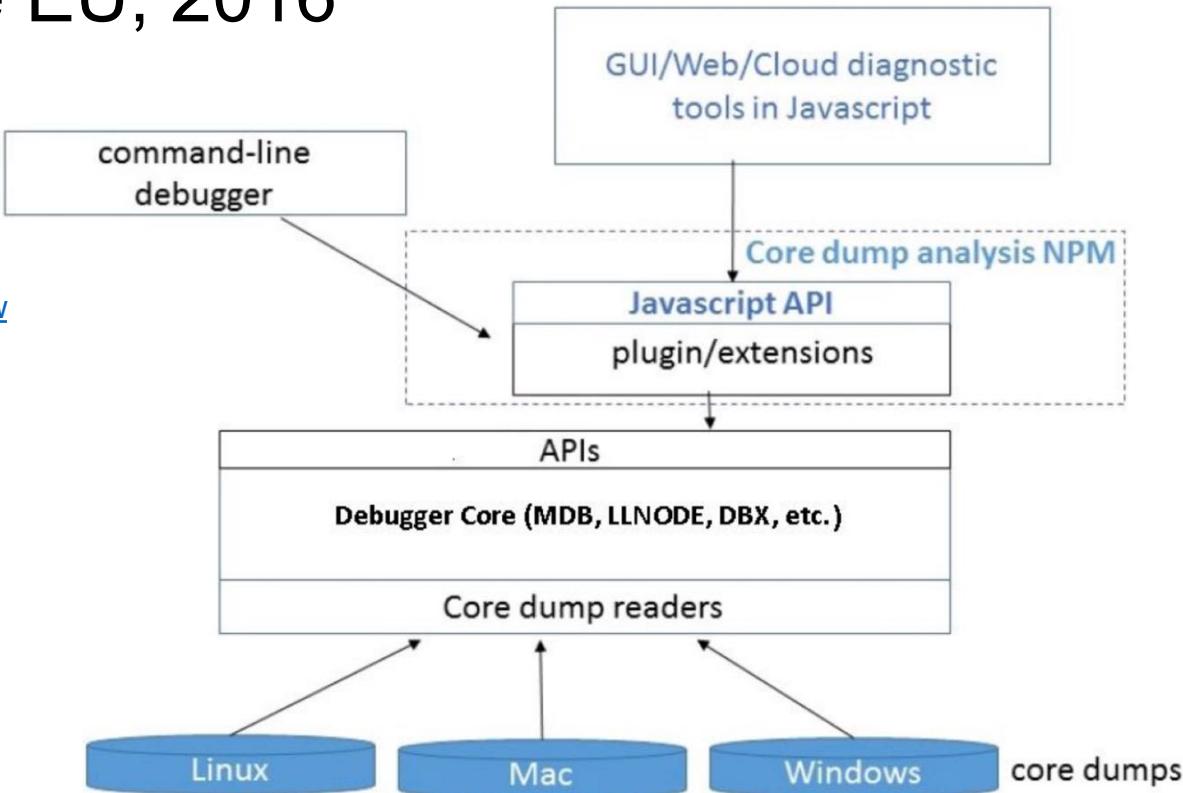


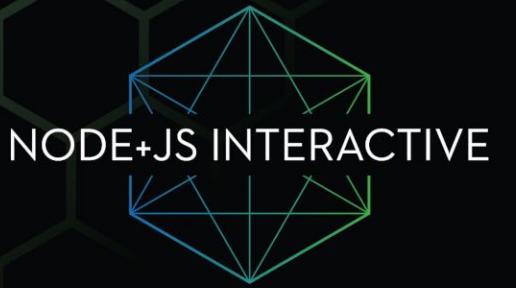
Future of Node/JS Post-mortem Diagnosis

Node Interactive EU, 2016

<https://www.slideshare.net/michaeldawson3572846/post-mortem-talk-node-interactive-eu>

<https://github.com/nodejs/diagnostics>





Thanks